

- **Name:** Iskander Nafikov
- **E-mail:** i.nafikov@innopolis.university
- **GitHub:** <https://github.com/iskanred/iu-devsecops-course/tree/main>
- **DockerHub:** <https://hub.docker.com/repository/docker/iskanred>
- **Username:** i.nafikov
- **Hostname:** macbook-KN70WX2PH

Task 1 - Preparation

1.1

Task description

Choose an application that has such characteristics as (points with "star" are not required but recommended):

- any non-confidential data like in JSON, XML or TOML format
- application configuration values (or environment variables)
- any secrets (like authorization credentials)
- availability from outside (on Internet) *
- has any health checks or/and status endpoints *

It could be any yours application or open-source project that meets the requirements. You are even able to work with different applications in different lab tasks but try to avoid it. However, it's **much better** to work with the single project but it's acceptable to involve only the required application functionality regarding the particular task requirements. Put the git link to your chosen project into report and provide a brief project description for what it is and how it works.

- For this lab I updated my application `current-time-server` from the previous lab. I added new `/secret` endpoint to satisfy the requirement to have any secrets.
- Below is the brief project description of this application:

The service provides information about current time in Europe/Moscow timezone and also shows the number of getting this information (number of visits). Moreover, it provides a hashed value of a secret string that is configured as an environment variable.

- You can check more detailed description in the `README.md` of the application.

- Here are the links:
 - **GitHub:** <https://github.com/iskanred/iu-devsecops-course/tree/main/lab-04/current-time-server/>
 - **DockerHub:** <https://hub.docker.com/repository/docker/iskanred/current-time-server/general>

1.2

Task description

Get familiar with Kubernetes (k8s) and concepts.



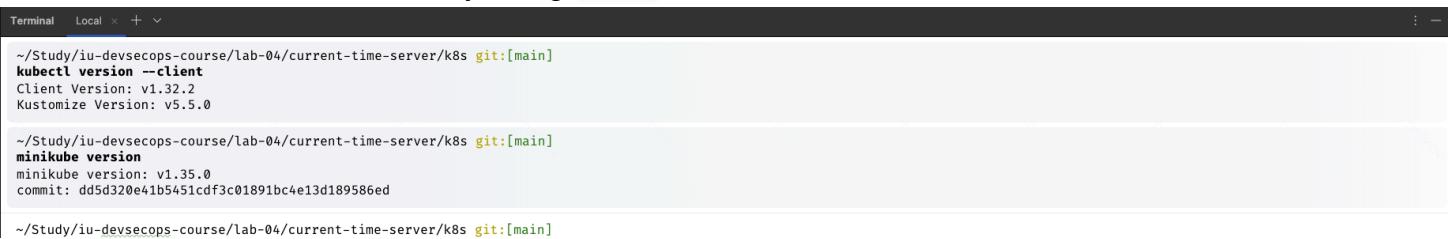
1.3

Task description

Install and set up the necessary tools:

- kubectl
- minikube or its alternative

- I installed these tools successfully using brew



```
Terminal Local × + × : - ~ /Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl version --client
Client Version: v1.32.2
Kustomize Version: v5.5.0

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
minikube version
minikube version: v1.35.0
commit: dd5d320e41b5451cdf3c01891bc4e13d189586ed

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
```

- What is more, I copied kube config values for minikube into a single `kubeconfig.yaml` file in a working directory

```

Local x kubeconfig.yaml
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat kubeconfig.yaml
apiVersion: v1

kind: Config
current-context: minikube

clusters:
- name: minikube
  cluster:
    server: https://127.0.0.1:32769
    certificate-authority: /Users/i.nafikov/.minikube/ca.crt
  extensions:
    - name: cluster_info
      extension:
        provider: minikube.sigs.k8s.io
        version: v1.35.0
        last-update: Tue, 18 Feb 2025 16:47:47 MSK

contexts:
- name: minikube
  context:
    cluster: minikube
    user: minikube
    namespace: default
  extensions:
    - name: context_info
      extension:
        provider: minikube.sigs.k8s.io
        version: v1.35.0
        last-update: Tue, 18 Feb 2025 16:47:47 MSK

users:
- name: minikube
  user:
    client-certificate: /Users/i.nafikov/.minikube/profiles/minikube/client.crt
    client-key: /Users/i.nafikov/.minikube/profiles/minikube/client.key

```

- Below is the content of this file

```

apiVersion: v1

kind: Config
current-context: minikube

clusters:
- name: minikube
  cluster:
    server: https://127.0.0.1:32769
    certificate-authority: /Users/i.nafikov/.minikube/ca.crt
  extensions:
    - name: cluster_info
      extension:
        provider: minikube.sigs.k8s.io
        version: v1.35.0
        last-update: Tue, 18 Feb 2025 16:47:47 MSK

contexts:
- name: minikube
  context:
    cluster: minikube
    user: minikube
    namespace: default
  extensions:
    - name: context_info
      extension:
        provider: minikube.sigs.k8s.io
        version: v1.35.0
        last-update: Tue, 18 Feb 2025 16:47:47 MSK

```

users:

- name: minikube
- user:
 - client-certificate: /Users/i.nafikov/.minikube/profiles/minikube/client.crt
 - client-key: /Users/i.nafikov/.minikube/profiles/minikube/client.key

- I can check if I copied properties correctly

```
kubectl --kubeconfig kubeconfig.yaml config get-contexts
```



```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl --kubeconfig kubeconfig.yaml config get-contexts
CURRENT      NAME        CLUSTER      AUTHINFO      NAMESPACE
*           minikube    minikube    minikube    default
```

- Seem so!

1.4

Task description

Get access to Kubernetes Dashboard.

- I accessed Kubernetes Dashboard using minikube

```
minikube dashboard
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
minikube dashboard
  ↗ Enabling dashboard ...
    - Using image docker.io/kubernetesui/dashboard:v2.7.0
    - Using image docker.io/kubernetesui/metrics-scraper:v1.0.8
💡 Some dashboard features require the metrics-server addon. To enable all features please run:

  minikube addons enable metrics-server

💡 Verifying dashboard health ...
  Launching proxy ...
  Verifying proxy health ...
  Opening http://127.0.0.1:63957/api/v1/namespaces/kubernetes-dashboard/services/http:kubernetes-dashboard:/proxy/ in your default browser ...
```

The screenshot shows the Kubernetes Dashboard interface. On the left is a sidebar with navigation links: Workloads, Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets, Service, Ingresses, Ingress Classes, Services, Config and Storage, Config Maps, Persistent Volume Claims, Secrets, Storage Classes, Cluster, Cluster Role Bindings, Cluster Roles, Events, Namespaces, Network Policies, and Nodes.

Workload Status:

- Daemon Sets: Running: 1
- Deployments: Running: 7
- Pods: Running: 13
- Replica Sets: Running: 7

Daemon Sets:

Name	Namespace	Images	Labels	Pods	Created
kube-proxy	kube-system	registry.k8s.io/kube-proxy:v1.32.0	k8s-app: kube-proxy	1 / 1	4 hours ago

Deployments:

Name	Namespace	Images	Labels	Pods	Created
dashboard-metrics-scraper	kubernetes-dashboard	docker.io/kubernetesui/metrics-scraper:v1.0@sha256:760a4987f07a0474dc334fc5d3569b9529bf985b2d2d9f356092ce206e98765c	addonmanager.kubernetes.io/mode: Reconcile k8s-app: dashboard-metrics-scraper kubernetes.io/minikube-addons: dashboard add-onmanager.kubernetes.io/mode: Reconcile	1 / 1	7 minutes ago
kubernetes-dashboard	kubernetes-dashboard	docker.io/kubernetesui/dashboard:v2.7.0@sha256:2e500d29e9d5f4a086b908ebbdfe7ecac57d2ab0965b24f588b1d449841ef93	k8s-app: kubernetes-dashboard kubernetes.io/minikube-addons: dashboard	1 / 1	7 minutes ago

- Also, it's better to mention in the task description how to access it because at first I tried to deploy it by my own using **helm** by [this](#) instruction but got into a trouble 😅

Task 2 - k8s Nodes

We are going to start to learn Kubernetes from `kubectl` command line that is dedicated to work with k8s cluster

1

✍ Task description

After installing all required lab tools and starting `minikube` cluster, use `kubectl` commands to `get` and `describe` your cluster nodes.

- Using `get` we can get list our nodes

```
kubectl get node
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get node
NAME      STATUS   ROLES      AGE     VERSION
minikube  Ready    control-plane  4h56m   v1.32.0
```

- We see that there is only one node in our cluster: `minikube` with a `control-plane` role
- Using `describe` we can get detailed information about our nodes

```
kubectl describe node minikube
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl describe node minikube
Name:           minikube
Roles:          control-plane
Labels:         beta.kubernetes.io/arch=arm64
                beta.kubernetes.io/os=linux
                kubernetes.io/arch=arm64
                kubernetes.io/hostname=minikube
                kubernetes.io/os=linux
Annotations:   minikube.k8s.io/commit=dd5d320e41b5451cdf3c01891bc4e13d189586ed
                minikube.k8s.io/name=minikube
                minikube.k8s.io/role=control-plane
                minikube.k8s.io/updated_at=2025_02_18T16_47_0700
                minikube.k8s.io/version=v1.35.0
                node-role.kubernetes.io/control-plane=
                node.kubernetes.io/exclude-from-external-load-balancers=
                kubeadm.alpha.kubernetes.io/cri-socket: unix:///var/run/cri-dockerd.sock
                node.alpha.kubernetes.io/ttl: 0
                volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Tue, 18 Feb 2025 16:47:44 +0300
Taints:          <none>
Unschedulable:   false
Labels:          holderIdentity: minikube
                 acquireTime: <unset>
                 renewTime: Tue, 18 Feb 2025 21:46:43 +0300
Conditions:     Type        Status  LastHeartbeatTime           LastTransitionTime        Reason           Message
                MemoryPressure False   Tue, 18 Feb 2025 21:45:53 +0300  Tue, 18 Feb 2025 16:47:43 +0300  KubeletHasSufficientMemory  Kubelet has sufficient memory available
                DiskPressure   False   Tue, 18 Feb 2025 21:45:53 +0300  Tue, 18 Feb 2025 16:47:43 +0300  KubeletHasNoDiskPressure  Kubelet has no disk pressure
                PIDPressure   False   Tue, 18 Feb 2025 21:45:53 +0300  Tue, 18 Feb 2025 16:47:43 +0300  KubeletHasSufficientPID   Kubelet has sufficient PID available
                Ready         True    Tue, 18 Feb 2025 21:45:53 +0300  Tue, 18 Feb 2025 16:47:45 +0300  KubeletReady            Kubelet is posting ready status
Addresses:
  InternalIP:  192.168.49.2
  Hostname:    minikube
Capacity:
  cpu:          4
  ephemeral-storage: 59848952Ki
  hugepages-1Gi: 0
  hugepages-2Mi: 0
  hugepages-32Mi: 0
  hugepages-64Ki: 0
  memory:       8113560Ki
  pods:          110
Allocatable:
  cpu:          4
  ephemeral-storage: 59848952Ki
  hugepages-1Gi: 0
  hugepages-2Mi: 0
  hugepages-32Mi: 0
  hugepages-64Ki: 0
  memory:       8113560Ki
  pods:          110
System Info:
  Machine ID:  f9ff83a2710f4d22a4f99feae8c5bdb4
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
```

- We see much information there: address of a node inside the Docker container, allocated resources, system info, k8s events, and etc.

2

Task description

Get the more detailed information about the particular node (it's fine if you have one node in the cluster).

- Basically the most detailed information can be obtained using `describe`
- However, we can also specify `wide` output for `get` command

```
kubectl get -o wide
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get node minikube -o wide
NAME      STATUS   ROLES      AGE     VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE        KERNEL-VERSION   CONTAINER-RUNTIME
minikube  Ready    control-plane  5h20m   v1.32.0   192.168.49.2  <none>        Ubuntu 22.04.5 LTS  6.8.0-31-generic  docker://27.4.1
```

3

Task description

Get the OS and CPU information.

- We can use `--output` or `-o` flag with an option `custom-columns` to specify what pieces of information we want to output

```
kubectl get node minikube -o=custom-  
columns='OS:.status.nodeInfo.osImage,CPU:.status.nodeInfo.architecture'
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl get node minikube -o=custom-columns='OS:.status.nodeInfo.osImage,CPU:.status.nodeInfo.architecture'  
OS CPU  
Ubuntu 22.04.5 LTS arm64
```

- Also, we can use `jsonpath` option

```
kubectl get node -o jsonpath='{range .items[*]}{"Name: '}{.metadata.name}{'\t'}  
{'OS: '}{.status.nodeInfo.osImage}{'\t'}{'CPU: '}{.status.nodeInfo.architecture}  
{'\n'}{end}'
```

```
Local Local 2 ...  
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl get nodes -o jsonpath='{range .items[*]}{"Name: '}{.metadata.name}{'\t'}{'OS: '}{.status.nodeInfo.osImage}{'\t'}{'CPU: '}{.status.nodeInfo.architecture}{'\n'}{end}'  
Name: minikube OS: Ubuntu 22.04.5 LTS CPU: arm64
```

4

Task description

Put the results into report.



Task 3 - k8s Pod

Pod is one of the simplest and basic k8s unit. It's like an abstraction over Docker containers. Inside pods containers run your application.

1

Task description

Figure out the necessary Pod spec fields

- **apiVersion** : This field specifies the API version of Kubernetes used.
- **kind** : This indicates the type of resource for creating. For pods it is `Pod`
- **metadata** : This section contains metadata about the `Pod`, such as its name, namespace, labels, and annotations.
- **spec** : This section contains the specification of the `Pod`, including the containers that should be created and ports that should be exposed to k8s virtual network.

2

Task description

Write a `Pod` spec for your chosen application, deploy the application and run a pod

- I created a file `pod.yaml` with `Pod` spec in my working directory

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
ls | grep pod
pod.yaml

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: current-time-server-pod
  labels:
    app: current-time-server
spec:
  containers:
    - name: current-time-server
      image: iskanred/current-time-server:1.0.0
      ports:
        - containerPort: 8080
```

- Here is the `pod.yaml`

```
apiVersion: v1
kind: Pod
metadata:
  name: current-time-server-pod
  labels:
    app: current-time-server
spec:
  containers:
    - name: current-time-server
      image: iskanred/current-time-server:1.0.0
      ports:
        - containerPort: 8080
```

3

Task description

With `kubectl`, get the pods, pod logs, describe pod, go into pod shell.

```
Local Local 2 pod.yaml :  
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl apply -f pod.yaml  
pod/current-time-server-pod created  
  
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl get po  
NAME READY STATUS RESTARTS AGE  
current-time-server-pod 0/1 ContainerCreating 0 4s  
  
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl get po  
NAME READY STATUS RESTARTS AGE  
current-time-server-pod 0/1 ContainerCreating 0 13s  
  
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl get po  
NAME READY STATUS RESTARTS AGE  
current-time-server-pod 1/1 Running 0 29s  
  
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
```

4

Task description

Make sure that your app is working correctly inside Pod

- Using `exec` command I executed a command inside a Pod's container

```
kubectl exec current-time-server-pod -- curl localhost:8080/time -w '\n'
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl exec current-time-server-pod -- curl localhost:8080/time -w '\n'  
 % Total    % Received % Xferd  Average Speed   Time     Time      Current  
          Dload Upload Total Spent   Left Speed  
100  8 100  8  0    0  572      0 --:--:-- --:--:-- --:--:-- 615  
22:57:36
```

- We see that the application works

5

Task description

Put the results into report



Task 4 - k8s Service

We use k8s Services to make an application accessible from outside the Kubernetes virtual network, provide an compatible IP address and link to DNS name to pod, to route internal/external traffic within pods.

1

Task description

Figure out the necessary `Service` spec fields

- `apiVersion` : This field specifies the API version of Kubernetes used.
- `kind` : This indicates the type of resource for creating. For services it is `Service`
- `metadata` : This section contains metadata about the `Service`, such as its name, namespace, labels, and annotations
- `spec` : This section contains the specification of the `Service`, including the service type, selector for pods
- `spec.ports` : This section defines the ports mapping configuration for the `Service`

2

Task description

Write a `Service` spec for your pod(s) and deploy the `Service`.

- First I change my `Pod` spec and created another one alongside:

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
ls | grep pod
pod-1.yaml
pod-2.yaml
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat pod-1.yaml
apiVersion: v1
kind: Pod
metadata:
  name: current-time-server-pod-1
  labels:
    app: current-time-server
spec:
  containers:
    - name: current-time-server
      image: iskanred/current-time-server:1.0.0
      ports:
        - containerPort: 8080
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat pod-2.yaml
apiVersion: v1
kind: Pod
metadata:
  name: current-time-server-pod-2
  labels:
    app: current-time-server
spec:
  containers:
    - name: current-time-server
      image: iskanred/current-time-server:1.0.0
      ports:
        - containerPort: 8080
```

- Then I created a Service using the following spec:

```

apiVersion: v1
kind: Service
metadata:
  name: current-time-server-service
spec:
  type: ClusterIP
  selector:
    app: current-time-server
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080

```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
ls | grep service
service.yaml
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat service.yaml
apiVersion: v1
kind: Service
metadata:
  name: current-time-server-service
spec:
  type: ClusterIP
  selector:
    app: current-time-server
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
```

- Finally, I deployed them all:

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f pod-1.yaml
pod/current-time-server-pod-1 created
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f pod-2.yaml
pod/current-time-server-pod-2 created
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f service.yaml
service/current-time-server-service created
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po
NAME          READY   STATUS    RESTARTS   AGE
current-time-server-pod-1  1/1     Running   0          15s
current-time-server-pod-2  1/1     Running   0          12s
```

3

Task description

With `kubectl`, get the Services and describe them.

- get

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get svc
NAME          TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
current-time-server-service  ClusterIP  10.105.84.72  <none>        80/TCP       105s
kubernetes     ClusterIP  10.96.0.1     <none>        443/TCP      6h50m
```

- **describe**

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl describe svc current-time-server-service
Name:           current-time-server-service
Namespace:      default
Labels:          <none>
Annotations:    <none>
Selector:       app=current-time-server
Type:           ClusterIP
IP Family Policy: SingleStack
IP Families:   IPv4
IP:             10.105.84.72
IPs:            10.105.84.72
Port:           <unset>  80/TCP
TargetPort:     8080/TCP
Endpoints:     10.244.0.23:8080,10.244.0.24:8080
Session Affinity: None
Internal Traffic Policy: Cluster
Events:         <none>
```

4

Task description

Make sure that pods can communicate between each other using DNS names, check pods addresses.

- We could see that the service has been actually deployed and explored my pods' endpoints

Properties				
Created	2m 52s ago	2025-02-18T23:36:16+03:00		
Name	current-time-server-service			
Namespace	default			
Annotations	endpoints.kubernetes.io/last-change-trigger-time=2025-02-18T20:36:16Z			
Subsets				
Addresses				
IP	Hostname	Target		
10.244.0.23		current-time-server-pod-1		
10.244.0.24		current-time-server-pod-2		
Ports				
Port	Name	Protocol		
8080		TCP		
Events				
No events found				

IP

- Here are the pods IP addresses inside the cluster

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po current-time-server-pod-1 -o=custom-columns='Pod IP:.status.podIP'
Pod IP
10.244.0.23
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po current-time-server-pod-2 -o=custom-columns='Pod IP:.status.podIP'
Pod IP
10.244.0.24
```

- We can make intercommunication without a service by these IP

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec current-time-server-pod-1 -- curl 10.244.0.24:8080/time -w '\n'
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
               Dload  Upload Total Spent   Left Speed
100     8 100      8    0      0  2466      0 --::-- --::-- --::-- 2666
23:43:22
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec current-time-server-pod-2 -- curl 10.244.0.23:8080/time -w '\n'
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
               Dload  Upload Total Spent   Left Speed
100     8 100      8    0      0  2064      0 --::-- --::-- --::-- 2666
23:43:11
```

- However, now it is also possible to communicate by the same IP through the service:

10.105.84.72 and 80 port

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec current-time-server-pod-2 -- curl 10.105.84.72/time -w '\n'
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
               Dload  Upload Total Spent   Left Speed
100     8 100      8    0      0  2418      0 --::-- --::-- --::-- 2666
23:47:19
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec current-time-server-pod-1 -- curl 10.105.84.72/time -w '\n'
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
               Dload  Upload Total Spent   Left Speed
100     8 100      8    0      0   908      0 --::-- --::-- --::-- 888
23:48:01
```

DNS

- Apart from that we can now refer to the service by DNS name

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec current-time-server-pod-1 -- curl http://current-time-server-service/time -w '\n'
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
               Dload  Upload Total Spent   Left Speed
100     8 100      8    0      0  1842      0 --::-- --::-- --::-- 2000
00:55:58
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec current-time-server-pod-2 -- curl http://current-time-server-service/time -w '\n'
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
               Dload  Upload Total Spent   Left Speed
100     8 100      8    0      0  1812      0 --::-- --::-- --::-- 2000
00:59:09
```

- This is possible because of the DNS server kybe-system/coredns and kybe-system/kube-dns service

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl -n kube-system get deploy
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
coredns   1/1     1           1           23m
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl -n kube-system get svc
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
kube-dns  ClusterIP  10.96.0.10  <none>        53/UDP,53/TCP,9153/TCP   23m
```

- We see that our pods have predefined domains that are resolved by the kube-dns service

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec -it current-time-server-pod-1 -- cat /etc/resolv.conf
nameserver 10.96.0.10
search default.svc.cluster.local svc.cluster.local cluster.local tcsbank.ru
options ndots:5

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec -it current-time-server-pod-2 -- cat /etc/resolv.conf
nameserver 10.96.0.10
search default.svc.cluster.local svc.cluster.local cluster.local tcsbank.ru
options ndots:5
```

5

Task description

Delete any Pod , recreate it and check addresses again. Make sure that traffic is routed to the new Pod correctly

- I deleted the pod current-time-server-pod-2 and create it again

```
kubectl delete -f pod-2.yaml
kubectl apply -f pod-2.yaml
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl delete -f pod-2.yaml
pod "current-time-server-pod-2" deleted

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -n default
NAME      READY   STATUS    RESTARTS   AGE
current-time-server-pod-1  1/1     Running   1 (13m ago)  20h

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f pod-2.yaml
pod/current-time-server-pod-2 created

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -n default
NAME      READY   STATUS    RESTARTS   AGE
current-time-server-pod-1  1/1     Running   1 (15m ago)  20h
current-time-server-pod-2  1/1     Running   0          27s
```

- We can notice that this pod changed its IP address

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po current-time-server-pod-2 --output=custom-columns='IP:.status.podIP'
IP
10.244.0.13
```

- Then I made an HTTP request to my service 10 times

```
for i in {1..10}; do kubectl exec current-time-server-pod-1 -- curl
http://current-time-server-service/time; done
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
for i in {1..10}; do kubectl exec current-time-server-pod-1 -- curl http://current-time-server-service/time; done
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
                                         Dload  Upload Total Spent   Left Speed
100     8  100     8    0    0      871      0  --::--  --::--  --::-- 888
21:31:05  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
                                         Dload  Upload Total Spent   Left Speed
100     8  100     8    0    0      2347      0  --::--  --::--  --::-- 2666
21:31:05  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
                                         Dload  Upload Total Spent   Left Speed
100     8  100     8    0    0      990      0  --::--  --::--  --::-- 1000
21:31:05  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
                                         Dload  Upload Total Spent   Left Speed
100     8  100     8    0    0      1863      0  --::--  --::--  --::-- 200021:31:05
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
                                         Dload  Upload Total Spent   Left Speed
100     8  100     8    0    0      2643      0  --::--  --::--  --::-- 400021:31:06
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
                                         Dload  Upload Total Spent   Left Speed
100     8  100     8    0    0      2255      0  --::--  --::--  --::-- 266621:31:06
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
                                         Dload  Upload Total Spent   Left Speed
100     8  100     8    0    0      2391      0  --::--  --::--  --::-- 266621:31:06
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
                                         Dload  Upload Total Spent   Left Speed
100     8  100     8    0    0      2889      0  --::--  --::--  --::-- 4000
21:31:06  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
                                         Dload  Upload Total Spent   Left Speed
100     8  100     8    0    0      2642      0  --::--  --::--  --::-- 400021:31:06
  % Total    % Received % Xferd  Average Speed   Time   Time   Time Current
                                         Dload  Upload Total Spent   Left Speed
100     8  100     8    0    0      2955      0  --::--  --::--  --::-- 400021:31:06
```

- We can see the this pods has corresponding logs

```
kubectl logs current-time-server-pod-2
```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]

Task description

Learn about `Loadbalancer` and `NodePort`.

ClusterIP

- **Description:** This is the default Service type. Such a Service exposes the service on a cluster-internal IP. It makes the Service accessible only from within the cluster.
- **Use Case:** Necessary for internal communication between Pods within the cluster. For example, microservices that need to communicate with each other can use a `ClusterIP` Service.

NodePort

- **Description:** A `NodePort` service exposes the `Service` on each Node's IP at a static port. A request to the Node's IP at the `NodePort` is forwarded to the `Service`.
- **Use Case:** Useful for development and testing, or for when you want to expose the `Service` externally without a load balancer. It's a straightforward method to expose a service available via any of the cluster nodes.

LoadBalancer

- **Description:** This type of Service provisions an external load balancer (if the cloud provider supports this feature), which directs traffic to the `ClusterIP` of the `Service`. The load balancer is assigned a public IP address.
- **Use Case:** Suitable for production applications where you want to expose a `Service` to the internet or external clients. This is commonly used with cloud providers like AWS, GCP, or Azure.

ExternalName

- **Description:** An `ExternalName` Service maps the `Service` to a DNS name. It does not proxy any traffic but serves as a way to reference an external service based on a name.
- **Use Case:** Useful for integrating external services into the Kubernetes cluster. Instead of using an IP or a fixed address, you can reference the external service by its DNS name.

7

Task description

Deploy an external `Service` to access your application from outside, e.g., from your local host

- I decided to use `LoadBalancer` service type
- Here is the updated `service.yaml` file content

```

apiVersion: v1
kind: Service
metadata:
  name: current-time-server-service
spec:
  type: LoadBalancer
  selector:
    app: current-time-server
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
  externalIPs:
    - 192.168.49.2

```

- Now let's recreate our service

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f service.yaml
service/current-time-server-service configured

```

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po current-time-server-service -o=wide
Error from server (NotFound): pods "current-time-server-service" not found

```

Exit code 1 ⓘ

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get svc current-time-server-service -o=wide
NAME          TYPE        CLUSTER-IP      EXTERNAL-IP     PORT(S)      AGE      SELECTOR
current-time-server-service   LoadBalancer   10.111.137.238  192.168.49.2   80:30467/TCP  21h     app=current-time-server

```

- First, let's try to connect to the minikube node that is running inside a Docker container

```

docker exec -it minikube curl http://localhost:30467/time -w '\n'

```

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
docker exec -it minikube curl http://localhost:30467/time -w '\n'
22:18:34

```

- After, let's I tunnel node's connection to my host machine

```

minikube service current-time-server-service

```

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
minikube service current-time-server-service

```

NAMESPACE	NAME	TARGET PORT	URL
default	current-time-server-service	80	http://192.168.49.2:30467

⚠ Starting tunnel for service current-time-server-service.

NAMESPACE	NAME	TARGET PORT	URL
default	current-time-server-service		http://127.0.0.1:52413

⚠ Opening service default/current-time-server-service in default browser...
! Because you are using a Docker driver on darwin, the terminal needs to be open to run it.

- Now we can access it locally



- This is funny to understand that there are so many ports in such a chain:

Pod's 8080 --> Node's Internal 80 --> Node's External 30467 --> Host's 52413

8

Task description

Put the results into report



Task 5 - k8s Deployment

Deployment is Kubernetes manifest to manage Pods automatically by [Controller](#). It helps to release, scale and upgrade Pods.

1

Task description

Figure out the necessary `Deployment` spec fields.

- **apiVersion** : This field specifies the API version of Kubernetes used.
- **kind** : This indicates the type of resource for creating. For deployments it is `Deployment`
- **metadata** : This section contains metadata about the Deployment, such as its name, namespace, labels, and annotations. The name of the Deployment is particularly important for managing and identifying it.
- **spec** : The specification for the Deployment, which includes the following necessary fields:
 - **replicas** : Specifies the desired number of Pod replicas. This is how many instances of your application you want to run.
 - **selector** : A label selector that identifies the Pods that belong to this Deployment. It must match the labels defined in the Pod template.
- **template** : A Pod template used to create the Pods. It contains the specification of the Pods, including metadata and spec for the containers.

2

Task description

Make sure that you wiped previous `Pod` manifests. Write a `Deployment` spec for your pod(s) and deploy the application.

- Firstly, I deleted pods

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl delete -f pod-1.yaml
pod "current-time-server-pod-1" deleted
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl delete -f pod-2.yaml
pod "current-time-server-pod-2" deleted
```

- I created `deployment.yaml` file with 2 pod replicas remaining pods' configuration the same

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
ls | grep deployment
deployment.yaml
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: current-time-server-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: current-time-server
  template:
    metadata:
      labels:
        app: current-time-server
  spec:
    containers:
      - name: current-time-server
        image: iskanred/current-time-server:1.0.0
        ports:
          - containerPort: 8080
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
```

- Below is the content of this file

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: current-time-server-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: current-time-server
  template:
    metadata:
      labels:
        app: current-time-server
  spec:
    containers:
      - name: current-time-server
        image: iskanred/current-time-server:1.0.0
```

ports:
– containerPort: 8080

- Here is the result of deploying the deployment

```
kubectl apply -f deployment.yaml
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f deployment.yaml
deployment.apps/current-time-server-deployment created

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get deploy -o=wide
NAME           READY   UP-TO-DATE   AVAILABLE   AGE   CONTAINERS   IMAGES   SELECTOR
current-time-server-deployment   2/2     2           2           66s   current-time-server   iskanred/current-time-server:1.0.0
app=current-time-server

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -o=wide
NAME          READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS
GATES
current-time-server-deployment-555c44f85d-k5btb   1/1     Running   0          91s   10.244.0.19   minikube   <none>
current-time-server-deployment-555c44f85d-p57g4   1/1     Running   0          91s   10.244.0.20   minikube   <none>
```

3

Task description

With `kubectl`, get the Deployments and describe them.

- I called get and describe commands for the deployment

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get deploy -o=wide
NAME           READY   UP-TO-DATE   AVAILABLE   AGE   CONTAINERS   IMAGES   SELECTOR
current-time-server-deployment   2/2     2           2           3m9s  current-time-server   iskanred/current-time-server:1.0.0
app=current-time-server

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl describe deploy
Name:           current-time-server-deployment
Namespace:      default
CreationTimestamp:  Wed, 19 Feb 2025 22:52:49 +0300
Labels:          <none>
Annotations:    deployment.kubernetes.io/revision: 1
Selector:        app=current-time-server
Replicas:       2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=current-time-server
  Containers:
    current-time-server:
      Image:      iskanred/current-time-server:1.0.0
      Port:       8080/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:    <none>
      Volumes:   <none>
      Node-Selectors: <none>
      Tolerations: <none>
  Conditions:
    Type        Status  Reason
    Available   True    MinimumReplicasAvailable
    Progressing  True    NewReplicaSetAvailable
OldReplicaSets: <none>
NewReplicaSet:  current-time-server-deployment-555c44f85d (2/2 replicas created)
Events:
  Type      Reason          Age      From            Message
  Normal   ScalingReplicaSet  3m19s   deployment-controller  Scaled up replica set current-time-server-deployment-555c44f85d from 0 to 2

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
```

- We here

- `StrategyType` for upgrading the deployment with its configuration
- `NewReplicaSet` for controlling pods' replicas creation and management
- and other useful information

4

Task description

Update your `Deployment` manifest to scale your application to three replicas.

- The most convenient way to scale the number of replicas of a running deployment is to use `kubectl scale` command.
- It may be useful for temporary increasing the number of pods in conditions of high traffic load.
- Here is the scaling to 3 replicas:

```
kubectl scale deployment current-time-server-deployment --replicas=3
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl scale deployment current-time-server-deployment --replicas=3
deployment.apps/current-time-server-deployment scaled
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get deploy
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
current-time-server-deployment   3/3     3           3           11m
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -o=wide
zsh: command not found: kubectl
```

Exit code 127 ⓘ

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -o=wide
NAME          READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS
GATES
current-time-server-deployment-555c44f85d-8dhv8   1/1   Running   0          25s   10.244.0.21   minikube   <none>   <none>
current-time-server-deployment-555c44f85d-k5btb   1/1   Running   0          11m   10.244.0.19   minikube   <none>   <none>
current-time-server-deployment-555c44f85d-p57g4   1/1   Running   0          11m   10.244.0.20   minikube   <none>   <none>
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
```

5

Task description

Access pod shell and logs using `Deployment` labels.

- We access all deployment's pods logs

```
kubectl logs -l app=current-time-server
```

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl logs -l app=currentTime-server
2025-02-19T20:04:18.324Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : Starting CurrentTimeServerApplicationKt v1.0.0 using Java 17.0.11 with PID 1
(/current-time-server/current-time-server-1.0.0.jar started by user in /current-time-server)
2025-02-19T20:04:18.325Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : No active profile set, falling back to 1 default profile: "default"
2025-02-19T20:04:19.338Z INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port(s): 8080 (http)
2025-02-19T20:04:19.343Z INFO 1 --- [main] o.apache.catalina.core.StandardService : Starting service [Tomcat]
2025-02-19T20:04:19.343Z INFO 1 --- [main] o.apache.catalina.core.StandardEngine : Starting Servlet engine: [Apache Tomcat/10.1.13]
2025-02-19T20:04:19.417Z INFO 1 --- [main] o.a.c.c.C.[Tomcat].[localhost].[] : Initializing Spring embedded WebApplicationContext
2025-02-19T20:04:19.894Z INFO 1 --- [main] w.s.c.ServletWebServerApplicationContext : Root WebApplicationContext: initialization completed in 1010 ms
2025-02-19T20:04:19.931Z INFO 1 --- [main] o.s.b.a.e.web.EndpointLinksResolver : Exposing 4 endpoint(s) beneath base path '/actuator'
2025-02-19T20:04:19.942Z INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat started on port(s): 8080 (http) with context path ''
2025-02-19T20:04:19.942Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : Started CurrentTimeServerApplicationKt in 1.928 seconds (process running for 2.202)
2025-02-19T19:52:51.425Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : Starting CurrentTimeServerApplicationKt v1.0.0 using Java 17.0.11 with PID 1
(/current-time-server/current-time-server-1.0.0.jar started by user in /current-time-server)
2025-02-19T19:52:51.426Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : No active profile set, falling back to 1 default profile: "default"
2025-02-19T19:52:53.694Z INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port(s): 8080 (http)
2025-02-19T19:52:53.702Z INFO 1 --- [main] o.apache.catalina.core.StandardService : Starting service [Tomcat]
2025-02-19T19:52:53.703Z INFO 1 --- [main] o.apache.catalina.core.StandardEngine : Starting Servlet engine: [Apache Tomcat/10.1.13]
2025-02-19T19:52:53.791Z INFO 1 --- [main] o.a.c.c.C.[Tomcat].[localhost].[] : Initializing Spring embedded WebApplicationContext
2025-02-19T19:52:53.792Z INFO 1 --- [main] w.s.c.ServletWebServerApplicationContext : Root WebApplicationContext: initialization completed in 2166 ms
2025-02-19T19:52:54.714Z INFO 1 --- [main] o.s.b.a.e.web.EndpointLinksResolver : Exposing 4 endpoint(s) beneath base path '/actuator'
2025-02-19T19:52:54.833Z INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat started on port(s): 8080 (http) with context path ''
2025-02-19T19:52:54.901Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : Started CurrentTimeServerApplicationKt in 3.989 seconds (process running for 4.758)
2025-02-19T19:52:51.307Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : Starting CurrentTimeServerApplicationKt v1.0.0 using Java 17.0.11 with PID 1
(/current-time-server/current-time-server-1.0.0.jar started by user in /current-time-server)
2025-02-19T19:52:51.309Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : No active profile set, falling back to 1 default profile: "default"
2025-02-19T19:52:53.611Z INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port(s): 8080 (http)
2025-02-19T19:52:53.617Z INFO 1 --- [main] o.apache.catalina.core.StandardService : Starting service [Tomcat]
2025-02-19T19:52:53.617Z INFO 1 --- [main] o.apache.catalina.core.StandardEngine : Starting Servlet engine: [Apache Tomcat/10.1.13]
2025-02-19T19:52:53.711Z INFO 1 --- [main] o.a.c.c.C.[Tomcat].[localhost].[] : Initializing Spring embedded WebApplicationContext
2025-02-19T19:52:53.712Z INFO 1 --- [main] w.s.c.ServletWebServerApplicationContext : Root WebApplicationContext: initialization completed in 2297 ms
2025-02-19T19:52:54.900Z INFO 1 --- [main] o.s.b.a.e.web.EndpointLinksResolver : Exposing 4 endpoint(s) beneath base path '/actuator'
2025-02-19T19:52:55.000Z INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat started on port(s): 8080 (http) with context path ''
2025-02-19T19:52:55.013Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : Started CurrentTimeServerApplicationKt in 4.378 seconds (process running for 4.869)

```

- Besides, for accessing logs of the specific we should list all the pods of the deployment using the label

```
kubectl get pods -l app=currentTime-server
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
```

```
kubectl get pods -l app=currentTime-server
```

NAME	READY	STATUS	RESTARTS	AGE
current-time-server-deployment-555c44f85d-8dhv8	1/1	Running	0	14m
current-time-server-deployment-555c44f85d-k5btb	1/1	Running	0	25m
current-time-server-deployment-555c44f85d-p57g4	1/1	Running	0	25m

- And then get logs of a desired pod

```
kubectl logs current-time-server-deployment-555c44f85d-8dhv8
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
```

```
kubectl logs current-time-server-deployment-555c44f85d-8dhv8
```



```

2025-02-19T20:04:18.324Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : Starting CurrentTimeServerApplicationKt v1.0.0 using Java 17.0.11 with PID 1 (/current-time-server/current-time-server-1.0.0.jar started by user in /current-time-server)
2025-02-19T20:04:18.325Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : No active profile set, falling back to 1 default profile: "default"
2025-02-19T20:04:19.338Z INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port(s): 8080 (http)
2025-02-19T20:04:19.343Z INFO 1 --- [main] o.apache.catalina.core.StandardService : Starting service [Tomcat]
2025-02-19T20:04:19.343Z INFO 1 --- [main] o.apache.catalina.core.StandardEngine : Starting Servlet engine: [Apache Tomcat/10.1.13]
2025-02-19T20:04:19.417Z INFO 1 --- [main] o.a.c.c.C.[Tomcat].[localhost].[] : Initializing Spring embedded WebApplicationContext
2025-02-19T20:04:19.417Z INFO 1 --- [main] w.s.c.ServletWebServerApplicationContext : Root WebApplicationContext: initialization completed in 1010 ms
2025-02-19T20:04:19.894Z INFO 1 --- [main] o.s.b.a.e.web.EndpointLinksResolver : Exposing 4 endpoint(s) beneath base path '/actuator'
2025-02-19T20:04:19.931Z INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat started on port(s): 8080 (http) with context path ''
2025-02-19T20:04:19.942Z INFO 1 --- [main] i.d.c.CurrentTimeServerApplicationKt : Started CurrentTimeServerApplicationKt in 1.928 seconds (process running for 2.202)

```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
```

- In the same way we can access pod's shell

```
kubectl exec -it current-time-server-deployment-555c44f85d-8dhv8 -- bash
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get pods -l app=current-time-server
NAME                   READY   STATUS    RESTARTS   AGE
current-time-server-deployment-555c44f85d-8dhv8   1/1     Running   0          20m
current-time-server-deployment-555c44f85d-k5btb   1/1     Running   0          31m
current-time-server-deployment-555c44f85d-p57g4   1/1     Running   0          31m
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec -it current-time-server-deployment-555c44f85d-8dhv8 -- bash
user@current-time-server-deployment-555c44f85d-8dhv8:/current-time-server$ exit
```

6

Task description

Make any application configuration change in your `Deployment` yaml and try to update the application. Monitor what are happened with pods (`--watch`).

- For application update I added a `command` for pods' containers

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: current-time-server-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: current-time-server
  template:
    metadata:
      labels:
        app: current-time-server
    spec:
      containers:
        - name: current-time-server
          image: iskanred/current-time-server:1.0.0
          command: ['sh', '-c', 'echo "Hello, Kubernetes!" && sleep 3600']
          ports:
            - containerPort: 8080
```

- I applied the new deployment configuration and started watching for pods

```
kubectl apply -f deployment.yaml && \
kubectl get pods -l app=current-time-server --watch
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat deployment.yaml | grep command
command: ['sh', '-c', 'echo "Hello, Kubernetes!" && sleep 3600']
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f deployment.yaml && kubectl get pods -l app=current-time-server --watch
deployment.apps/current-time-server-deployment configured
NAME          READY   STATUS    RESTARTS   AGE
current-time-server-deployment-555c44f85d-b9rzx  1/1    Running   0          20s
current-time-server-deployment-555c44f85d-mrfmk  1/1    Running   0          21s
current-time-server-deployment-8595ff49b6-4n5pf  1/1    Terminating   0          85s
current-time-server-deployment-8595ff49b6-jmszj  1/1    Terminating   0          84s
current-time-server-deployment-8595ff49b6-sncm2  0/1    ContainerCreating   0          0s
current-time-server-deployment-8595ff49b6-sncm2  1/1    Running   0          2s
current-time-server-deployment-555c44f85d-b9rzx  1/1    Terminating   0          22s
current-time-server-deployment-8595ff49b6-7zs7d  0/1    Pending    0          0s
current-time-server-deployment-8595ff49b6-7zs7d  0/1    Pending    0          0s
current-time-server-deployment-8595ff49b6-7zs7d  0/1    ContainerCreating   0          0s
current-time-server-deployment-555c44f85d-b9rzx  0/1    Error     0          22s
current-time-server-deployment-555c44f85d-b9rzx  0/1    Error     0          23s
current-time-server-deployment-555c44f85d-b9rzx  0/1    Error     0          23s
current-time-server-deployment-8595ff49b6-7zs7d  1/1    Running   0          1s
current-time-server-deployment-555c44f85d-mrfmk  1/1    Terminating   0          24s
current-time-server-deployment-555c44f85d-mrfmk  0/1    Error     0          24s
current-time-server-deployment-555c44f85d-mrfmk  0/1    Error     0          25s
current-time-server-deployment-555c44f85d-mrfmk  0/1    Error     0          25s
current-time-server-deployment-8595ff49b6-jmszj  0/1    Error     0          95s
current-time-server-deployment-8595ff49b6-jmszj  0/1    Error     0          95s
current-time-server-deployment-8595ff49b6-jmszj  0/1    Error     0          95s
current-time-server-deployment-8595ff49b6-4n5pf  0/1    Error     0          97s
current-time-server-deployment-8595ff49b6-4n5pf  0/1    Error     0          97s
current-time-server-deployment-8595ff49b6-4n5pf  0/1    Error     0          97s
```

- We can check if the application was updated by accessing its logs

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl logs -l app=current-time-server
Hello, Kubernetes!
Hello, Kubernetes!
```

7

Task description

Rollback to previous application version using `Deployment`.

- At first, let's check the history of the deployment

```
kubectl rollout history deployment/current-time-server-deployment
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl rollout history deployment/current-time-server-deployment
deployment.apps/current-time-server-deployment
REVISION  CHANGE-CAUSE
5        <none>
6        <none>
```

- We can notice that the `REVISION` is so hight because I debugged the deployment upgrading process in the previous task
- Now let's rollout the deployment to the previous revision

```
kubectl rollout undo deployment/current-time-server-deployment && \
kubectl get pods -l app=current-time-server --watch
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl rollout undo deployment/current-time-server-deployment && kubectl get pods -l app=current-time-server --watch
deployment.apps/current-time-server-deployment rolled back
NAME                      READY   STATUS        RESTARTS   AGE
current-time-server-deployment-555c44f85d-ltfkf   0/1    ContainerCreating   0          0s
current-time-server-deployment-8595ff49b6-7zs7d   1/1    Running       0          11m
current-time-server-deployment-8595ff49b6-sncm2   1/1    Running       0          11m
current-time-server-deployment-555c44f85d-ltfkf   1/1    Running       0          1s
current-time-server-deployment-8595ff49b6-7zs7d   1/1    Terminating   0          11m
current-time-server-deployment-555c44f85d-z9c7v   0/1    Pending       0          0s
current-time-server-deployment-555c44f85d-z9c7v   0/1    Pending       0          0s
current-time-server-deployment-555c44f85d-z9c7v   0/1    ContainerCreating   0          0s
current-time-server-deployment-555c44f85d-z9c7v   1/1    Running       0          1s
current-time-server-deployment-8595ff49b6-sncm2   1/1    Terminating   0          11m
^C
```

Exit code 1 ⚙

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl rollout status deployment/current-time-server-deployment
deployment "current-time-server-deployment" successfully rolled out
```

- We see the status of our rollout is successful
- Let's check pods' logs

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl logs -l app=current-time-server
2025-02-19T21:01:28.016Z INFO 1 --- [           main] i.d.c.CurrentTimeServerApplicationKt : Starting CurrentTimeServerApplicationKt v1.0.0 using Java 17.0.11 with PID 1 (/current-time-server/current-time-server-1.0.0.jar started by user in /current-time-server)
2025-02-19T21:01:28.017Z INFO 1 --- [           main] i.d.c.CurrentTimeServerApplicationKt : No active profile set, falling back to 1 default profile: "default"
2025-02-19T21:01:30.324Z INFO 1 --- [           main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port(s): 8080 (http)
2025-02-19T21:01:30.329Z INFO 1 --- [           main] o.apache.catalina.core.StandardService : Starting service [Tomcat]
2025-02-19T21:01:30.329Z INFO 1 --- [           main] o.apache.catalina.core.StandardEngine : Starting Servlet engine: [Apache Tomcat/10.1.13]
2025-02-19T21:01:30.432Z INFO 1 --- [           main] o.a.c.c.C.[Tomcat].[localhost].[] : Initializing Spring embedded WebApplicationContext
2025-02-19T21:01:30.434Z INFO 1 --- [           main] w.s.c.ServletWebServerApplicationContext : Root WebApplicationContext: initialization completed in 2321 ms
2025-02-19T21:01:31.521Z INFO 1 --- [           main] o.s.b.a.e.web.EndpointLinksResolver : Exposing 4 endpoint(s) beneath base path '/actuator'
2025-02-19T21:01:31.627Z INFO 1 --- [           main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat started on port(s): 8080 (http)
with context path ''
2025-02-19T21:01:31.691Z INFO 1 --- [4.183 seconds (process running for 4.51)
2025-02-19T21:01:29.205Z INFO 1 --- [           main] i.d.c.CurrentTimeServerApplicationKt : Starting CurrentTimeServerApplicationKt v1.0.0 using Java 17.0.11 with PID 1 (/current-time-server/current-time-server-1.0.0.jar started by user in /current-time-server)
2025-02-19T21:01:29.210Z INFO 1 --- [           main] i.d.c.CurrentTimeServerApplicationKt : No active profile set, falling back to 1 default profile: "default"
2025-02-19T21:01:31.592Z INFO 1 --- [           main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port(s): 8080 (http)
2025-02-19T21:01:31.599Z INFO 1 --- [           main] o.apache.catalina.core.StandardService : Starting service [Tomcat]
2025-02-19T21:01:31.599Z INFO 1 --- [           main] o.apache.catalina.core.StandardEngine : Starting Servlet engine: [Apache Tomcat/10.1.13]
2025-02-19T21:01:31.692Z INFO 1 --- [           main] o.a.c.c.C.[Tomcat].[localhost].[] : Initializing Spring embedded WebApplicationContext
2025-02-19T21:01:31.693Z INFO 1 --- [           main] w.s.c.ServletWebServerApplicationContext : Root WebApplicationContext: initialization completed in 2362 ms
2025-02-19T21:01:32.299Z INFO 1 --- [           main] o.s.b.a.e.web.EndpointLinksResolver : Exposing 4 endpoint(s) beneath base path '/actuator'
2025-02-19T21:01:32.345Z INFO 1 --- [           main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat started on port(s): 8080 (http)
with context path ''
2025-02-19T21:01:32.398Z INFO 1 --- [3.898 seconds (process running for 4.625)
```

- We see that the rollout was successful indeed!

- Finally, we can notice that our ReplicaSet s were updated by deployment controller

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl describe deploy
Name:           current-time-server-deployment
Namespace:      default
CreationTimestamp: Wed, 19 Feb 2025 22:52:49 +0300
Labels:         <none>
Annotations:   deployment.kubernetes.io/revision: 7
Selector:       app=current-time-server
Replicas:      2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType:  RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=current-time-server
  Containers:
    current-time-server:
      Image:      iskanred/current-time-server:1.0.0
      Port:       8080/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:    <none>
      Volumes:   <none>
      Node-Selectors: <none>
      Tolerations:  <none>
  Conditions:
    Type        Status  Reason
    Available   True    MinimumReplicasAvailable
    Progressing  True    NewReplicaSetAvailable
OldReplicaSets: current-time-server-deployment-8595ff49b6 (0/0 replicas created)
NewReplicaSet:  current-time-server-deployment-555c44f85d (2/2 replicas created)
Events:
  Type  Reason          Age   From            Message
  Normal ScalingReplicaSet 34m  deployment-controller  Scaled down replica set current-time-server-deployment-555c44f85d from 3 to 2
  Normal ScalingReplicaSet 33m  deployment-controller  (combined from similar events): Scaled up replica set current-time-server-deployment-555c44f85d from 0 to 1
  Normal ScalingReplicaSet 33m  deployment-controller  Scaled up replica set current-time-server-deployment-8595ff49b6 from 0 to 1
  Normal ScalingReplicaSet 33m  deployment-controller  Scaled down replica set current-time-server-deployment-555c44f85d from 2 to 1
  Normal ScalingReplicaSet 33m  deployment-controller  Scaled up replica set current-time-server-deployment-8595ff49b6 from 1 to 2
  Normal ScalingReplicaSet 33m  deployment-controller  Scaled down replica set current-time-server-deployment-555c44f85d from 1 to 0
  Normal ScalingReplicaSet 21m  deployment-controller  Scaled up replica set current-time-server-deployment-555c44f85d from 0 to 1
  Normal ScalingReplicaSet 21m  deployment-controller  Scaled down replica set current-time-server-deployment-8595ff49b6 from 2 to 1
  Normal ScalingReplicaSet 21m  deployment-controller  Scaled up replica set current-time-server-deployment-555c44f85d from 1 to 2
  Normal ScalingReplicaSet 71m  deployment-controller  Scaled down replica set current-time-server-deployment-8595ff49b6 from 1 to 0
```

8

Task description

Set up requests and limits for CPU and Memory for your application pods. Provide a PoC that it works properly. Explain and show what is happened when app reaches the CPU usage limit? And memory limit?

Exploring consumption

- Below is the updated deployment.yaml file. Notice that I removed container.command from the manifest

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: current-time-server-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: current-time-server
  template:
    metadata:
      labels:
```

```

app: current-time-server
spec:
  containers:
    - name: current-time-server
      image: iskanred/current-time-server:1.0.0
      ports:
        - containerPort: 8080
      resources:
        requests:
          memory: 256Mi
          cpu: 0.5
        limits:
          memory: 512Mi
          cpu: 1

```

- I applied this manifest

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: current-time-server-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: current-time-server
  template:
    metadata:
      labels:
        app: current-time-server
    spec:
      containers:
        - name: current-time-server
          image: iskanred/current-time-server:1.0.0
          ports:
            - containerPort: 8080
          resources:
            requests:
              memory: 256Mi
              cpu: 0.5
            limits:
              memory: 512Mi
              cpu: 1

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f deployment.yaml
deployment.apps/current-time-server-deployment configured

```

- We can see that these limits/requests were applied to our pods

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -o=custom-columns='Resources:.spec.containers[0].resources'
Resources
map[limits:map[cpu:1 memory:512Mi] requests:map[cpu:500m memory:256Mi]]
map[limits:map[cpu:1 memory:512Mi] requests:map[cpu:500m memory:256Mi]]

```

- What is more we can notice that their QoS is indeed Burstable because they do have memory and CPU limits/requests but they do not equal

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -o=custom-columns='Name:.metadata.name,QoS:.status.qosClass'
Name QoS
current-time-server-deployment-bd9664c9c-kdhfr Burstable
current-time-server-deployment-bd9664c9c-xtxlk Burstable

```

- Now let's check pods' metrics using minikube's metrics-server addon

```
minikube addons enable metrics-server
```

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
minikube addons list | grep metrics
| metrics-server | minikube | enabled ✓ | Kubernetes |

```

- Now we can monitor our pods' performance Uusing kubectl top command

```
kubectl top pods -l app=current-time-server
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl top pods -l app=current-time-server
NAME                  CPU(cores)   MEMORY(bytes)
current-time-server-deployment-bd9664c9c-kdhfr   2m          125Mi
current-time-server-deployment-bd9664c9c-xtxlk   2m          122Mi
```

- We see that my pods consume only 2m which is 0.002 of CPU core cycle time and about 125Mi of memory

Proof of concept

- We can infer that current thresholds for limits for both memory and CPU are too big since the application does no use so much
- Let's determine some other boundaries and test if they work actually!

CPU limits

My hypothesis is that **with smaller limit of CPU time my pods start slower.**

- This might happen because pods have lack of performance at the startup while application is configured: Spring beans should be created, JIT-compilation should "warm up" code regions, and etc.

Preparation

- To test startup duration I added `startupProbe` which access application's healthcheck endpoint

```
apiVersion: apps/v1
kind: Deployment
<...>
  containers:
    startupProbe:
      httpGet:
        path: /actuator/health
        port: 8080
      periodSeconds: 1
      failureThreshold: 10000
```

- We see that the prober is configured to check if a pod's container is started 10000 times every 1 second
- To measure how long the pod is starting I implemented a simple `measure.sh` script

```
#!/bin/zsh

PODS=$(kubectl get po -l app=current-time-server)

# Deleting pods if exist
if [[ -n $PODS ]]; then
```

```

echo "There are some pods with the same labels running"
kubectl delete -f deployment.yaml

while [[ $PODS != "" ]]
do
    PODS=$(kubectl get po -l app=current-time-server)
    echo "Waiting for the pods to be deleted..."
    sleep 1
done

echo "The pods are deleted"
fi

kubectl apply -f deployment.yaml
echo "Waiting for pods to start..."

seconds=0
while [[ $(kubectl get po -l app=current-time-server -o
jsonpath='{.items[0].status.containerStatuses[0].started}') == "false" ]]
do
    (( seconds+=1 ))
    sleep 1;
done

echo "Elapsed time: $seconds seconds"

```

- It removes already existed pods with desired labels and starts a deployment. After the deployment is started the script increments the number of seconds until the status is not `false`. Finally, the script prints the number of elapsed time for pod to start.
- What is more, to make the process of testing faster I decreased our deployment to `1` replica only

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: current-time-server-deployment
spec:
  replicas: 1
<...>

```

- In addition, I configured `requests = limits` to simplify the process of testing `limits` making pods to have Guaranteed QoS.
- Finally, I gave enough memory for my pods to be sure the memory is not a bottleneck when testing CPU.

- The final deployment.yaml is the following:

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: current-time-server-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: current-time-server
  template:
    metadata:
      labels:
        app: current-time-server
    spec:
      containers:
        - name: current-time-server
          image: iskanred/current-time-server:1.0.0
          ports:
            - containerPort: 8080
          resources:
            requests:
              memory: 1Gi
              cpu: 1
            limits:
              memory: 1Gi
              cpu: 1
      startupProbe:
        httpGet:
          path: /actuator/health
          port: 8080
        periodSeconds: 1
        failureThreshold: 10000
```

Test

- Let's check if our CPU limits bounds actually work
- Firstly, let's start with limits.cpu = 1

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat deployment.yaml | grep -B 2 cpu
  requests:
    memory: 1Gi
    cpu: 1
  limits:
    memory: 1Gi
    cpu: 1

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
./measure.sh
There are some pods with the same labels running
deployment.apps "current-time-server-deployment" deleted
Waiting for the pods to be deleted...
No resources found in default namespace.
Waiting for the pods to be deleted...
The pods are deleted
deployment.apps/current-time-server-deployment created
Waiting for pods to start...
Elapsed time: 5 seconds
```

- Now, let's change the value to limits.cpu = 0.5

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat deployment.yaml | grep -B 2 cpu
  requests:
    memory: 1Gi
    cpu: 0.5
  limits:
    memory: 1Gi
    cpu: 0.5

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
./measure.sh
There are some pods with the same labels running
deployment.apps "current-time-server-deployment" deleted
Waiting for the pods to be deleted...
No resources found in default namespace.
Waiting for the pods to be deleted...
The pods are deleted
deployment.apps/current-time-server-deployment created
Waiting for pods to start...
Elapsed time: 9 seconds
```

- Let's continue with `limits.cpu = 0.25`

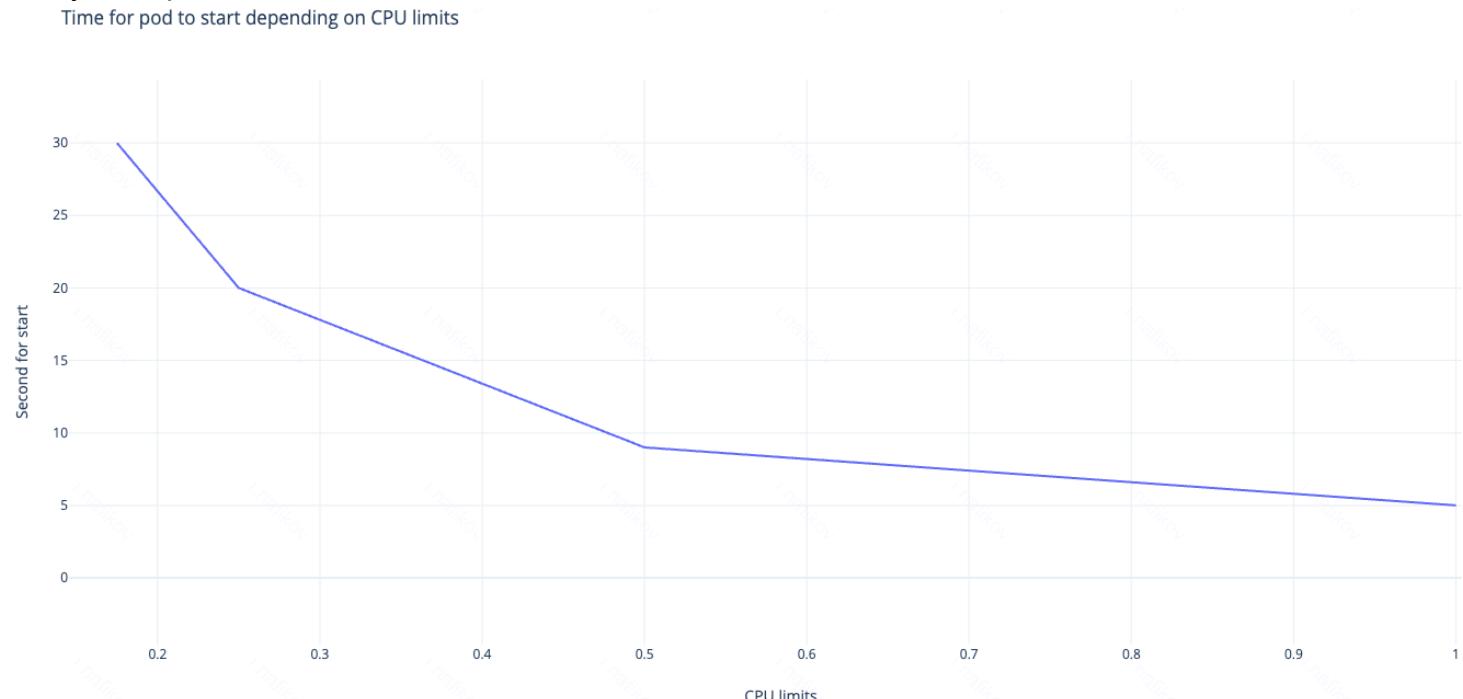
```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
./measure.sh
There are some pods with the same labels running
deployment.apps "current-time-server-deployment" deleted
Waiting for the pods to be deleted...
No resources found in default namespace.
Waiting for the pods to be deleted...
The pods are deleted
deployment.apps/current-time-server-deployment created
Waiting for pods to start...
Elapsed time: 20 seconds
```

- Finally, let's check the result with `limits.cpu = 0.175`

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat deployment.yaml | grep -B 2 cpu
  requests:
    memory: 1Gi
    cpu: 0.175
  limits:
    memory: 1Gi
    cpu: 0.175
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
./measure.sh
There are some pods with the same labels running
deployment.apps "current-time-server-deployment" deleted
Waiting for the pods to be deleted...
No resources found in default namespace.
Waiting for the pods to be deleted...
The pods are deleted
deployment.apps/current-time-server-deployment created
Waiting for pods to start...
Elapsed time: 30 seconds
```

- Okay, let's plot our results



- We see that **limits mechanism actually works!**. The smaller the `limits.cpu` the higher the amount of time for pod to start

Memory limits

My hypothesis is that **with smaller limit of memory pods may be killed due to OOM**

- This might happen if an application has not enough memory to start / work.

Preparation

- To check if `limits.memory` actually works the strategy is similar to check the `limits.cpu` on a startup.

- However, here we have a **binary** evaluation: either the pod will start successfully or not.
- So here I can just set `limit.memory` value so small that pods will be killed almost immediately.
- To make startup faster I set `limits.cpu = 1`.
- What is more, I kept `requests = limits`.

Test

- Let's check the value `limits.memory = 100Mi`

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl delete -f deployment.yaml
deployment.apps "current-time-server-deployment" deleted

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat deployment.yaml | grep -B 2 cpu
  requests:
    memory: 100Mi
    cpu: 1
  limits:
    memory: 100Mi
    cpu: 1

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f deployment.yaml
deployment.apps/current-time-server-deployment created

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get pod
NAME                  READY   STATUS      RESTARTS   AGE
current-time-server-deployment-67d6658b7d-9mjbt   0/1     CrashLoopBackOff   4 (58s ago)   2m40s

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -o jsonpath-as-json='{.items[0].status.containerStatuses[0].lastState}'
[
  {
    "terminated": {
      "containerID": "docker://b07e871b8f74bc1f97dcf9fcaebe60737e505ba54414450249262decf178da02",
      "exitCode": 137,
      "finishedAt": "2025-02-22T15:10:49Z",
      "reason": "OOMKilled",
      "startedAt": "2025-02-22T15:10:46Z"
    }
  }
]

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
```

- We see that our **limit.memory mechanism works too!** The container couldn't be started and tries to restart infinitely because kubelet kills it since it consumes more memory than the limit set.

9

Task description

Put the results into report



Task 6 - k8s Secrets

Secret manifest is a quite similar to configMap . However, we use Secret to work with confidential application data. Kubernetes encode secrets in Base64 format.

1

Task description

Figure out the necessary Secret spec fields.

- **apiVersion** : This field specifies the API version of Kubernetes used.
- **kind** : This indicates the type of resource for creating. For secrets it is `Secret`.
- **metadata** : This section contains metadata about the `Secret`, such as its name, namespace, labels, and annotations.
- **type** : Defines the type of Secret. Common types include:
 - `Opaque` : The default type for arbitrary data.
 - `kubernetes.io/dockerconfigjson` : For Docker registry credentials.
 - `kubernetes.io/basic-auth` : For basic authentication credentials.
 - `kubernetes.io/ssh-auth` : For SSH authentication key.
- **data** : The data fields where each entry is base64-encoded.
- **stringData** : This is similar to `data`, but it allows you to specify the values in plaintext. Kubernetes will encode them in base64 automatically when you create the Secret. This can be more convenient for users.
- It's worth to notice that secrets do not have `spec` field in its manifest.

2

Task description

Create and apply a new `Secret` manifest. For example, it could be login and password to login to your app or something else

- I created a secret from the following `secret.yaml` manifest:

```
apiVersion: v1
kind: Secret
metadata:
  name: current-time-server-secret
  labels:
    app: current-time-server
type: Opaque
stringData:
  secret: my-super-secret
```

- And applied it

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat secret.yaml
apiVersion: v1
kind: Secret
metadata:
  name: current-time-server-secret
  labels:
    app: current-time-server
type: Opaque
stringData:
  secret: my-super-secret

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f secret.yaml
secret/current-time-server-secret created

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get secret
NAME              TYPE      DATA   AGE
current-time-server-secret  Opaque   1      5s
```

3

Task description

With `kubectl` , get and describe your secret(s).

- get

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get secret -o wide
NAME              TYPE      DATA   AGE
current-time-server-secret  Opaque   1      105s

kubectl get secret current-time-server-secret -o json
{
  "apiVersion": "v1",
  "data": {
    "secret": "bXktc3VwZXItc2VjcmV0"
  },
  "kind": "Secret",
  "metadata": {
    "annotations": {
      "kubernetes.io/last-applied-configuration": "{\"apiVersion\":\"v1\",\"kind\":\"Secret\", \"metadata\":{\"annotations\":{},\"labels\":{\"app\":\"current-time-server\"},\"name\":\"current-time-server-secret\", \"namespace\":\"default\"},\"stringData\":{\"secret\":\"my-super-secret\"},\"type\":\"Opaque\"}\n"
    },
    "creationTimestamp": "2025-02-22T16:10:53Z",
    "labels": {
      "app": "current-time-server"
    },
    "name": "current-time-server-secret",
    "namespace": "default",
    "resourceVersion": "99157",
    "uid": "f32293ca-82be-4e45-b80f-f91aeb8ba804"
  },
  "type": "Opaque"
}
```

- describe

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl describe secret
Name:         current-time-server-secret
Namespace:    default
Labels:       app=current-time-server
Annotations: <none>

Type:  Opaque

Data
=====
secret: 15 bytes
```

4

Task description

Decode your secret(s).

- I decoded the value of my secret using base64

```
kubectl get secret current-time-server-secret -o jsonpath='{.data.secret}' |  
base64 -d
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl get secret current-time-server-secret -o jsonpath='{.data.secret}' | base64 -d && echo '\n'  
my-super-secret
```

5

Task description

Update your Deployment to reference to your secret as environment variable.

- I update the deployment.yaml as following:

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: current-time-server-deployment  
spec:  
  replicas: 2  
  selector:  
    matchLabels:  
      app: current-time-server  
  template:  
    metadata:  
      labels:  
        app: current-time-server  
    spec:  
      containers:  
        - name: current-time-server  
          image: iskanred/current-time-server:1.0.0  
          ports:  
            - containerPort: 8080  
      resources:  
        requests:  
          memory: 512Mi  
          cpu: 1  
      limits:
```

```

memory: 512Mi
cpu: 1
startupProbe:
  httpGet:
    path: /actuator/health
    port: 8080
  periodSeconds: 1
  failureThreshold: 10000
env:
  - name: SECRET
    valueFrom:
      secretKeyRef:
        name: current-time-server-secret
        key: secret

```

- Then I applied it

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat deployment.yaml | grep -A 5 env
```

```

env:
  - name: SECRET
    valueFrom:
      secretKeyRef:
        name: current-time-server-secret
        key: secret

```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f deployment.yaml
deployment.apps/current-time-server-deployment created
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -l app=current-time-server
NAME                   READY   STATUS    RESTARTS   AGE
current-time-server-deployment-5b469bc57-7pmvm  1/1     Running   0          18s
current-time-server-deployment-5b469bc57-hs9k9  1/1     Running   0          18s
```

6

Task description

Make sure that you are able to see your secret inside pod.

- To check if I can see my secret I tunneled the service to my local host using minikube

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get svc current-time-server-service
NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
current-time-server-service  LoadBalancer  10.98.173.1    192.168.49.2  80:30954/TCP  46h
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
minikube service current-time-server-service
```

NAMESPACE	NAME	TARGET PORT	URL
default	current-time-server-service	80	http://192.168.49.2:30954

Starting tunnel for service current-time-server-service.

NAMESPACE	NAME	TARGET PORT	URL
default	current-time-server-service		http://127.0.0.1:53950

Opening service default/current-time-server-service in default browser...
! Because you are using a Docker driver on darwin, the terminal needs to be open to run it.■

- Finally, I could easily get the **MD5**-hashed value of my secret by accessing the `/secret` endpoint of my application

```
~/Study/iu-devsecops-course/lab-04/current-time-server git:[main]
curl http://127.0.0.1:53950/secret -w '\n'
MD5 hash of secret is: ad44f1be904b9375641b666bd4bbc531
```

- The response is

MD5 hash of secret is: ad44f1be904b9375641b666bd4bbc531

- Let's compare this value with the hash value of my secret:

```
~/Study/iu-devsecops-course/lab-04/current-time-server git:[main]
curl http://127.0.0.1:53950/secret | awk '{print $6}'
% Total    % Received % Xferd  Average Speed   Time     Time      Current
          Dload  Upload Total Spent   Left Speed
100  55 100  55  0    0  8587  0 --:-- --:-- --:-- 9166
ad44f1be904b9375641b666bd4bbc531
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server git:[main]
echo -n "my-super-secret" | md5
ad44f1be904b9375641b666bd4bbc531
```

- We see that they are actually equal

7

Task description

Put the results into report



Task 7 - k8s configMap

ConfigMap is Kubernetes manifest to store application configuration setting in two ways: key-value pairs as environment variables and text (usually JSON) data as dedicated file into container filesystem. We usually use configMap to store non sensitive data as plain text

1

Task description

Figure out the necessary `configMap` spec fields.

- apiVersion** : This field specifies the API version of Kubernetes used.
- kind** : This indicates the type of resource for creating. For config maps it is `ConfigMap`.
- metadata** : This section contains metadata about the `ConfigMap`, such as its name, namespace, labels, and annotations.

- **data** : A map of key-value pairs, where the keys are strings and the values can be strings or binary data encoded in base64. This field holds the configuration information.
- **binaryData** : (optional) Similar to `data`, but here values must be base64 encoded. It's used for binary data (e.g., a configuration file stored in binary format).

2

Task description

Modify your `Deployment` manifest to set up some app configuration via environment variables

- I decided to use the same value for obtaining a secret but now from a config map, not a secret
- I reused the same `deployment.yaml` but edited the secret source to config map

```
<...>
env:
  - name: SECRET
    valueFrom:
      configMapKeyRef:
        name: current-time-server-config
        key: secret
```

3

Task description

Create a new `configMap` manifest. In data spec, put some app configuration as key-value pair (it could be the same as in previous exercise). In the `Deployment.Pod` spec add the connection to key-value pair from `configMap` yaml file

- Below is my `configmap.yaml`

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: current-time-server-config
  labels:
    app: current-time-server
data:
  secret: my-config-secret
```

- I applied it

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f configmap.yaml
configmap/current-time-server-config created

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get configmap -l app=current-time-server
NAME          DATA   AGE
current-time-server-config   1      21s

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get configmap -l app=current-time-server -o json
{
  "apiVersion": "v1",
  "items": [
    {
      "apiVersion": "v1",
      "data": {
        "secret": "my-config-secret"
      },
      "kind": "ConfigMap",
      "metadata": {
        "annotations": {
          "kubectl.kubernetes.io/last-applied-configuration": "{\"apiVersion\":\"v1\", \"data\":{\"secret\":\"my-config-secret\"},\"kind\":\"ConfigMap\",\"metadata\":{},\"labels\":{\"app\":\"current-time-server\"},\"name\":\"current-time-server-config\",\"namespace\":\"default\"}\n",
          "creationTimestamp": "2025-02-22T17:14:57Z",
          "labels": {
            "app": "current-time-server"
          },
          "name": "current-time-server-config",
          "namespace": "default",
          "resourceVersion": "102338",
          "uid": "2a9c1d04-06af-4562-b263-1a1b59110b08"
        }
      },
      "kind": "List",
      "metadata": {
        "resourceVersion": ""
      }
    }
  ],
  "kind": "ConfigMap"
}
```

- Then, I applied my deployment.yaml

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl delete -f secret.yaml
secret "current-time-server-secret" deleted

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f deployment.yaml
deployment.apps/current-time-server-deployment configured
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -l app=current-time-server
NAME           READY   STATUS    RESTARTS   AGE
current-time-server-deployment-bbd4db585-hdd6w  1/1     Running   0          30s
current-time-server-deployment-bbd4db585-pmlpz  1/1     Running   0          25s
```

- Let's check again if secret exist inside my pods

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get svc current-time-server-service
NAME          TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
current-time-server-service   LoadBalancer   10.98.173.1   192.168.49.2   80:30954/TCP   47h

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
minikube service current-time-server-service


| NAMESPACE | NAME                        | TARGET PORT | URL                                                               |
|-----------|-----------------------------|-------------|-------------------------------------------------------------------|
| default   | current-time-server-service | 80          | <a href="http://192.168.49.2:30954">http://192.168.49.2:30954</a> |


Starting tunnel for service current-time-server-service.


| NAMESPACE | NAME                        | TARGET PORT | URL                                                         |
|-----------|-----------------------------|-------------|-------------------------------------------------------------|
| default   | current-time-server-service |             | <a href="http://127.0.0.1:56853">http://127.0.0.1:56853</a> |


Opening service default/current-time-server-service in default browser...
! Because you are using a Docker driver on darwin, the terminal needs to be open to run it.
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server git:[main]
curl http://127.0.0.1:56853/secret | awk '{print $6}'f
% Total    % Received % Xferd  Average Speed   Time     Time      Time  Current
          Dload Upload Total Spent   Left Speed
100  55  100  55    0      0  5126      0 --:--:-- --:--:-- --:--:--  5500
e8d4a4f2a01da1d337b45d94b67c7687
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server git:[main]
echo -n "my-config-secret" | md5
e8d4a4f2a01da1d337b45d94b67c7687
```

- We see that it worked!

4

Task description

Create a new file like `config.json` file and put some json data into

- Here is `config.json`:

```
{
  "application": "current-time-server",
  "object": {
    "name": "value"
  }
}
```

5

Task description

Create a new one `configMap` manifest. Connect `configMap` yaml file with `config.json` file to read the data from it.

- I created the following `current-time-server-config-json`:

```
kubectl create configmap current-time-server-config-json --from-file config.json
```

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl create configmap current-time-server-config-json --from-file=config.json
configmap/current-time-server-config-json created

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl label configmap current-time-server-config-json app=current-time-server --overwrite=true
configmap/current-time-server-config-json labeled

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get configmap current-time-server-config-json -o=yaml
apiVersion: v1
data:
  config.json: |
    {
      "application": "current-time-server",
      "object": {
        "name": "value"
      }
    }
kind: ConfigMap
metadata:
  creationTimestamp: "2025-02-22T17:40:22Z"
  labels:
    app: current-time-server
  name: current-time-server-config-json
  namespace: default
  resourceVersion: "103656"
  uid: 8acd2e11-d01b-40e8-b505-755bf20a4da7

```

6

Task description

Update your Deployment to add Volumes and VolumeMounts .

- I updated my deployment.yaml

```

< . . . >

volumeMounts:
  - name: current-time-server-config-json-volume
    mountPath: /current-time-server

volumes:
  - name: current-time-server-config-json-volume
    configMap:
      name: current-time-server-config-json

```

- And applied it

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
cat deployment.yaml | grep -B 10 -A 10 volume
  path: /actuator/health
  port: 8080
  periodSeconds: 1
  failureThreshold: 10000
env:
- name: SECRET
  valueFrom:
    configMapKeyRef:
      name: current-time-server-config
      key: secret
volumeMounts:
- name: current-time-server-config-json-volume
  mountPath: /etc/config
volumes:
- name: current-time-server-config-json-volume
  configMap:
    name: current-time-server-config-json
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f deployment.yaml
deployment.apps/current-time-server-deployment created
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -l app=current-time-server
NAME                      READY   STATUS    RESTARTS   AGE
current-time-server-deployment-685779988b-lktkv   1/1     Running   0          8s
current-time-server-deployment-685779988b-nbx5p   1/1     Running   0          8s
```

7

Task description

With `kubectl`, check the `configMap` details. Make sure that you see the data as plain text

- We can check it see that it displays as a JSON text:

```
kubectl get configmap current-time-server-config-json -o
jsonpath='{.data.config\.json}'
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get configmap current-time-server-config-json -o jsonpath='{.data.config\.json}'
{
  "application": "current-time-server",
  "object": {
    "name": "value"
  }
}
```

- However it also displays as a plain text when we do not specify config field

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get configmap current-time-server-config-json -o jsonpath='{.data}'
{"config.json": "{\n  \"application\": \"current-time-server\",\n  \"object\": {\n    \"name\": \"value\"\n  }\n}"}
```

8

Task description

Check the filesystem inside app container to show the loaded file data on the specified path.

- Let's check pods' filesystem

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po
NAME                      READY   STATUS    RESTARTS   AGE
current-time-server-deployment-685779988b-lktkv   1/1     Running   0          7m28s
current-time-server-deployment-685779988b-nbx5p   1/1     Running   0          7m28s

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec current-time-server-deployment-685779988b-lktkv -- ls /etc/config
config.json

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec current-time-server-deployment-685779988b-lktkv -- cat /etc/config/config.json
{
  "application": "current-time-server",
  "object": {
    "name": "value"
  }
}

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec current-time-server-deployment-685779988b-nbx5p -- ls /etc/config
config.json

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec current-time-server-deployment-685779988b-nbx5p -- cat /etc/config/config.json
{
  "application": "current-time-server",
  "object": {
    "name": "value"
  }
}
```

- And we see that the volume `current-time-server-config-json-volume` were actually mounted to both my pods!

9

Task description

Put the results into report



Task 8 - k8s Namespace

Namespace Kubernetes Manifest is designed for different projects and deployment environments isolation. With Namespaces we can separate App 1 deployment from App 2 deployment, manage (and isolate) cluster resources for them, define users list to have access either to App 1 or to App 2 deployment. Using Namespaces , we also can define a different environments like DEV, TEST, STAGE. In that way, Namespaces is a required feature for a real Kubernetes production clusters.

1

Task description

Figure out the necessary `Namespace` spec fields.

- `apiVersion` : This field specifies the API version of Kubernetes used.
- `kind` : This indicates the type of resource for creating. For namespace it is `Namespace` .
- `metadata` : This section contains metadata about the `ConfigMap` , such as its name, namespace, labels, and annotations.
 - `name` : The name of the Namespace. This must be a unique identifier **within the cluster** and typically follows DNS label formatting requirements (lowercase alphanumeric characters or '-' and must start and end with an alphanumeric character).

2

Task description

Create a two different `Namespaces` in your k8s cluster

- I created two manifests: `nginx-namespace.yaml` and `current-time-server.yaml`
- `nginx-namespace.yaml`

```
apiVersion: v1
kind: Namespace
metadata:
  name: nginx-namespace
  labels:
    app: nginx
```

- `current-time-server.yaml`

```
apiVersion: v1
kind: Namespace
metadata:
  name: current-time-server-namespace
  labels:
    app: current-time-server
```

- And applied them

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f nginx-namespace.yaml
namespace/nginx-namespace created
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f current-time-server-namespace.yaml
namespace/current-time-server-namespace created
```

3

Task description

Using `kubectl` , get and describe your Namespaces .

- `get`

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get namespace -l app -o=wide
NAME           STATUS   AGE
current-time-server-namespace  Active   74s
nginx-namespace          Active   82s
```

- `describe`

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl describe namespace -l app
Name:         current-time-server-namespace
Labels:       app=current-time-server
              kubernetes.io/metadata.name=current-time-server-namespace
Annotations: <none>
Status:      Active

No resource quota.

No LimitRange resource.

Name:         nginx-namespace
Labels:       app=nginx
              kubernetes.io/metadata.name=nginx-namespace
Annotations: <none>
Status:      Active

No resource quota.

No LimitRange resource.
```

4

Task description

Deploy two different applications in two different Namespaces with `kubectl` . By the way, it's acceptable even just to deploy the same objects (same previous app) in the different Namespaces but with different resources names.

- It may be already clear that I am going to deploy my `current-time-server` application with all its resources to `current-time-server-namespace` and a simple `nginx` application to `nginx-namespace` .

current-time-server

- I created all the necessary resources

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl create configmap current-time-server-config-json --from-file config.json -n current-time-server-namespace
configmap/current-time-server-config-json created

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f secret.yaml -n current-time-server-namespace
secret/current-time-server-secret created

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f service.yaml -n current-time-server-namespace
service/current-time-server-service created

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl apply -f deployment.yaml -n current-time-server-namespace
deployment.apps/current-time-server-deployment created
```

- And it worked

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -n current-time-server-namespace
NAME                      READY   STATUS    RESTARTS   AGE
current-time-server-deployment-698d7d6ff8-qb9z2   1/1     Running   0          3m21s
current-time-server-deployment-698d7d6ff8-tlcgq   1/1     Running   0          3m26s

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get svc -n current-time-server-namespace
NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
current-time-server-service   LoadBalancer   10.101.210.172  192.168.49.2   80:32254/TCP   5m36s

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
docker exec minikube curl -s http://localhost:32254/time
21:34:36
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
docker exec minikube curl -s http://localhost:32254/secret
MD5 hash of secret is: ad44f1be904b9375641b666bd4bbc531
```

nginx

- I created a simple Deployment with an **nginx** container. The manifest for it is `nginx-deployment.yaml`:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
```

```
  ports:  
    - containerPort: 80
```

- Also, I created a simple Service for make this Deployment accessible from outside. The manifest for it is `nginx-service.yaml`:

```
apiVersion: v1  
kind: Service  
metadata:  
  name: nginx-service  
spec:  
  type: LoadBalancer  
  selector:  
    app: nginx  
  ports:  
    - protocol: TCP  
      port: 80  
      targetPort: 80  
  externalIPs:  
    - 192.168.49.2
```

- I applied these resources inside the `nginx-namespace` namespace

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl apply -f nginx-deployment.yaml -n nginx-namespace  
deployment.apps/nginx-deployment created  
  
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl apply -f nginx-service.yaml -n nginx-namespace  
service/nginx-service created  
  
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl get po -n nginx-namespace  
NAME           READY   STATUS    RESTARTS   AGE  
nginx-deployment-96b9d695-mfl7c   1/1     Running   0          22s  
  
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
kubectl get svc -n nginx-namespace  
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE  
nginx-service   LoadBalancer   10.102.218.176   192.168.49.2   80:32212/TCP   13s
```

- Finally, I could access the **nginx** application

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]  
docker exec minikube curl -s http://localhost:32212/ -w '\n'  
<!DOCTYPE html>  
<html>  
<head>  
<title>Welcome to nginx!</title>  
<style>  
html { color-scheme: light dark; }  
body { width: 35em; margin: 0 auto;  
font-family: Tahoma, Verdana, Arial, sans-serif; }  
</style>  
</head>  
<body>  
<h1>Welcome to nginx!</h1>  
<p>If you see this page, the nginx web server is successfully installed and  
working. Further configuration is required.</p>  
  
<p>For online documentation and support please refer to  
<a href="http://nginx.org/">nginx.org</a>.<br/>  
Commercial support is available at  
<a href="http://nginx.com/">nginx.com</a>.</p>  
  
<p><em>Thank you for using nginx.</em></p>  
</body>  
</html>
```

Task description

With `kubectl`, get and describe pods from different Namespaces with `-n` flag.

- **get**

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -n current-time-server-namespace
NAME           READY   STATUS    RESTARTS   AGE
current-time-server-deployment-698d7d6ff8-qb9z2  1/1     Running   0          14m
current-time-server-deployment-698d7d6ff8-tlcgq   1/1     Running   0          14m
```

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -n nginx-namespace
NAME           READY   STATUS    RESTARTS   AGE
nginx-deployment-96b9d695-mfl7c   1/1     Running   0          5m55s
```

- **describe**

```
kubectl describe po -n current-time-server-namespace
Name:         current-time-server-deployment-698d7d6ff8-qb9z2
Namespace:    current-time-server-namespace
Priority:    0
Service Account: default
Node:        minikube/192.168.49.2
Start Time:   Sat, 22 Feb 2025 21:31:08 +0300
Labels:       app=current-time-server
              pod-template-hash=698d7d6ff8
Annotations: <none>
Status:      Running
IP:          10.244.0.120
IPs:
  IP:        10.244.0.120
Controlled By: ReplicaSet/current-time-server-deployment-698d7d6ff8
Containers:
  current-time-server:
    Container ID: docker://fee4b3f8161429c6c68cb48daef146539040d1413a810794e3b50917afcf366e
    Image:        iskanred/current-time-server:1.0.0
    Image ID:    docker-pullable://iskanred/current-time-server@sha256:b78f5bec42d299048c0bde0c6e6a853db3ac555cc888b491dc90b367b9da5a3d
    Port:        8080/TCP
    Host Port:  0/TCP
    State:      Running
      Started:  Sat, 22 Feb 2025 21:31:08 +0300
    Ready:      True
    Restart Count: 0
    Limits:
      cpu:  1
      memory: 512Mi
    Requests:
      cpu:  1
      memory: 512Mi
    Startup:  http-get http://:8080/actuator/health delay=0s timeout=1s period=1s #success=1 #failure=10000
    Environment:
      SECRET: <set to the key 'secret' in secret 'current-time-server-secret'> Optional: false
    Mounts:
      /etc/config from current-time-server-config-json-volume (rw)
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-48gj9 (ro)
Conditions:
  Type        Status
  PodReadyToStartContainers  True
  Initialized  True
  Ready        True
  ContainersReady  True
```

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl describe po -n nginx-namespace
Name:           nginx-deployment-96b9d695-mfl7c
Namespace:      nginx-namespace
Priority:       0
Service Account: default
Node:           minikube/192.168.49.2
Start Time:     Sat, 22 Feb 2025 21:39:49 +0300
Labels:         app=nginx
                pod-template-hash=96b9d695
Annotations:    <none>
Status:         Running
IP:             10.244.0.121
IPs:
  IP:          10.244.0.121
Controlled By: ReplicaSet/nginx-deployment-96b9d695
Containers:
  nginx:
    Container ID: docker://42ac554acbba9364aa9b9c5cab4c4a8361de8ee61eba7acbd0c936f893a1fcf
    Image:          nginx:latest
    Image ID:      docker-pullable://nginx@sha256:91734281c0ebfc6f1aea979cffeed5079cf786228a71cc6f1f46a228cde6e34
    Port:          80/TCP
    Host Port:    0/TCP
    State:        Running
      Started:   Sat, 22 Feb 2025 21:40:06 +0300
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-2pkjz (ro)
Conditions:
  Type        Status
  PodReadyToStartContainers  True
  Initialized  True
  Ready        True
  ContainersReady  True
  PodScheduled  True
Volumes:
  kube-api-access-2pkjz:
    Type:          Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:  kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI:   true
    Options:       <nil>
QoS Class:  BestEffort

```

6

Task description

Can you see and can you connect to the resources from different Namespaces?

- Yes. For example in nginx-namespace

```

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -n nginx-namespace
NAME            READY   STATUS    RESTARTS   AGE
nginx-deployment-96b9d695-mfl7c  1/1     Running   0          8m41s

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec -n nginx-namespace nginx-deployment-96b9d695-mfl7c -- curl -s http://localhost:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>

```

- Or in current-time-server-namespace

```
~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl get po -n current-time-server-namespace
NAME                   READY   STATUS    RESTARTS   AGE
current-time-server-deployment-698d7d6ff8-qb9z2   1/1     Running   0          18m
current-time-server-deployment-698d7d6ff8-tlcgq   1/1     Running   0          18m

~/Study/iu-devsecops-course/lab-04/current-time-server/k8s git:[main]
kubectl exec -n current-time-server-namespace current-time-server-deployment-698d7d6ff8-qb9z2 -- curl -s http://localhost:8080/time -w '\n'
21:49:34
```

7

Task description

Put the results into report



Post scriptum

- You can find all the materials of this lab in my GitHub repo of the course
<https://github.com/iskanred/iu-devsecops-course/tree/main/lab-04>
- Here is my file structure with files related to k8s for this lab



